

AI Deep Learning: Feedforward Neural Networks

Thuan L Nguyen, PhD

Slide 2: AI Deep Learning: Feedforward Neural Networks (FFNN)



AI Deep learning (Source: mindovermachines.com)

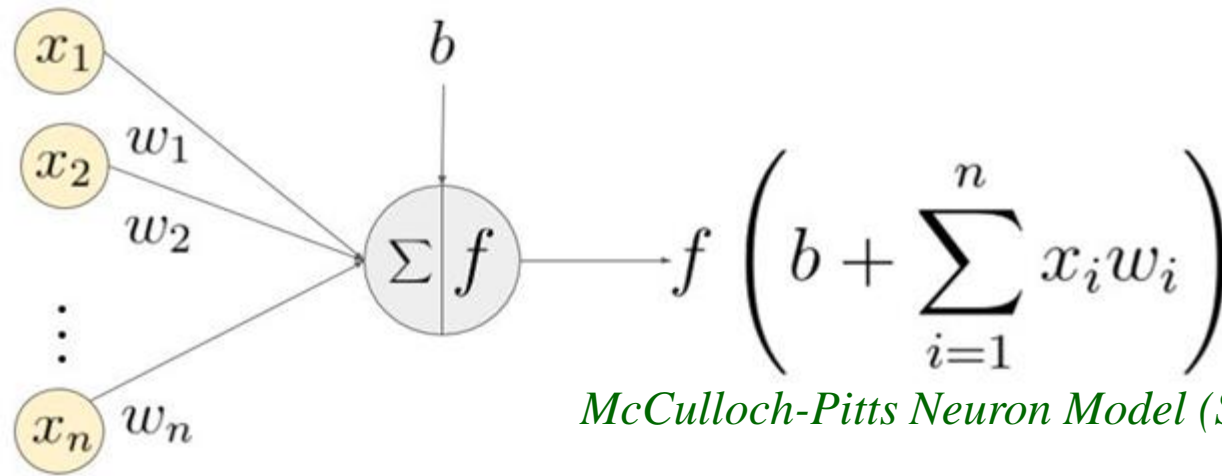
Slide 3: AI Deep Learning: Feedforward Neural Networks (FFNN)

1. Feedforward Neural Networks: Neurons and Perceptron
2. Feedforward Neural Networks: Overview
3. Feedforward Neural Networks: Architecture
4. Feedforward Neural Networks: Gradient-Based Learning
5. Feedforward Neural Networks: Optimization & Cost Function
6. Feedforward Neural Networks: Backpropagation Algorithm

Slide 4: AI Deep Learning: Feedforward Neural Networks (FFNN)

Artificial Neurons: The McCulloch-Pitts Neuron

A **simple rate coding model** of real neurons is also known as a **Threshold Logic Unit** :

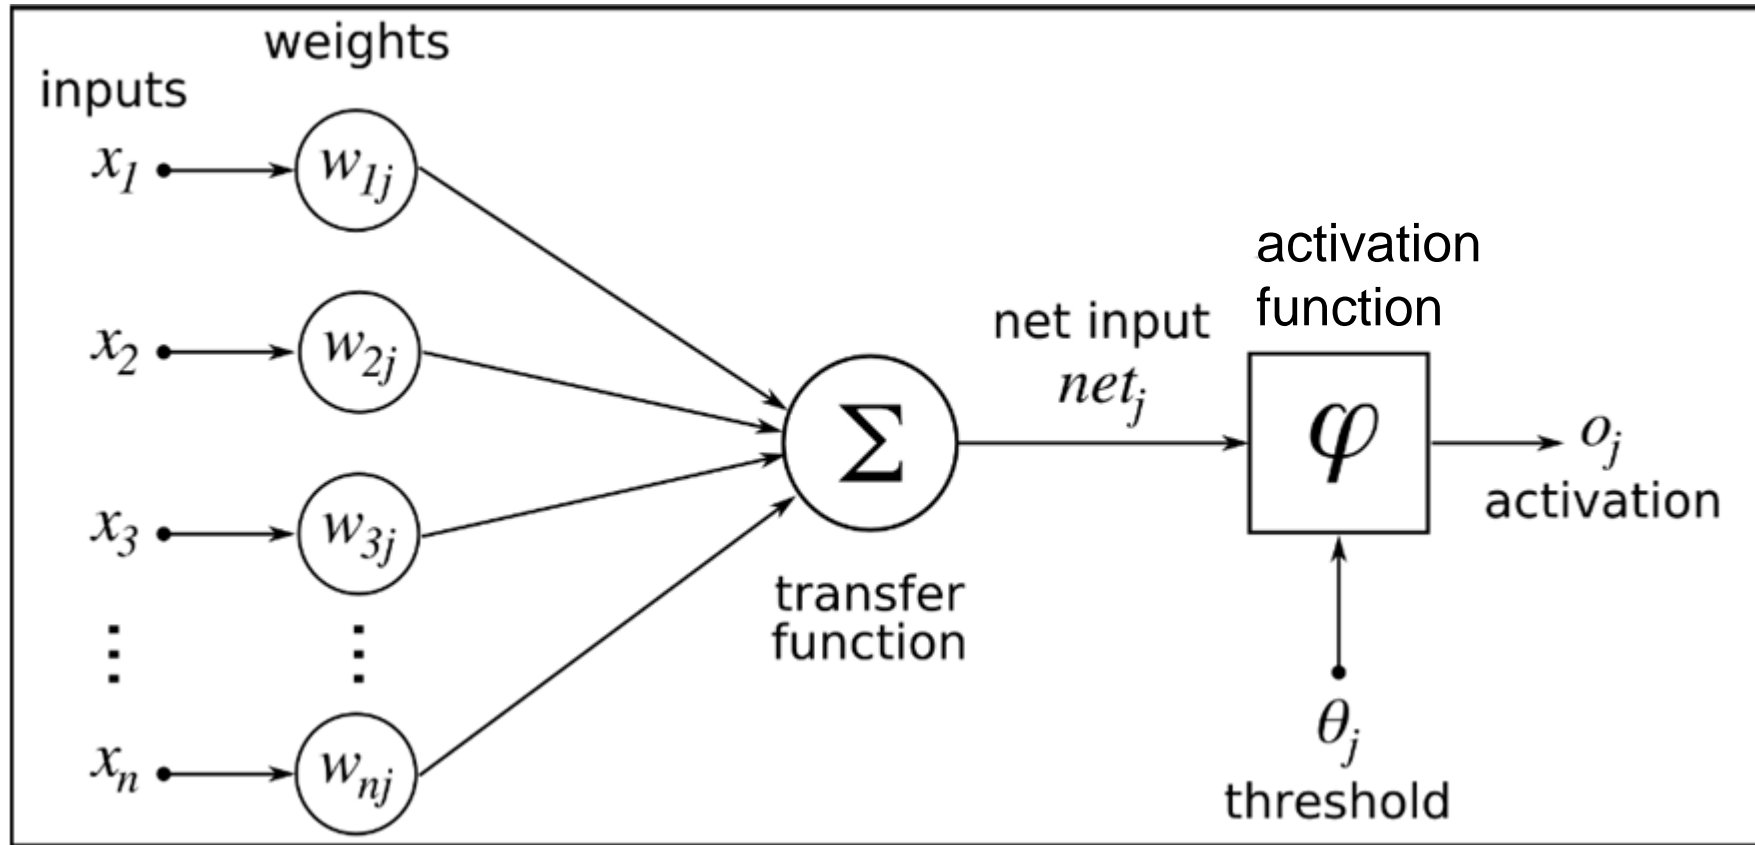


McCulloch-Pitts Neuron Model (Source:Wikipedia.com)

- A set of **synapses** (i.e. connections) brings in **activations, i.e., inputs**, from other neurons.
- A **processing unit** sums the inputs, and then applies a **non-linear activation function**
 - Is also often called a threshold or transfer or squashing function
- An **output line** transmits the result to other neurons.

Slide 5: AI Deep Learning: Feedforward Neural Networks (FFNN)

Perceptron: A Network of The McCulloch-Pitts Neurons




McCulloch-Pitts Neuron model (Sources: wikipedia.org)

Slide 6: AI Deep Learning: Feedforward Neural Networks (FFNN)

Deep Learning: Simple Single-Layer Neural Networks

Perceptron:

- The fundamental unit of an artificial neural network
- A simple – single-layer – artificial neural network:
 - A simple neural network that has one layer of input neurons feeding forward to one output layer of McCulloch-Pitts neurons, with full connectivity.



Perceptron: Single-Layer Neural Network (Sources: Wikipedia)

Slide 7: AI Deep Learning: Feedforward Neural Networks (FFNN)

Perceptron: A Network of The McCulloch-Pitts Neurons

- Frank Rosenblatt introduced the concept of perceptron (1958)
 - Each input I_i is multiplied by a weight w_{ji} (synaptic strength)
 - These weighted inputs are summed to give the activation level, A_j
 - The activation level is then transformed by an activation function to produce the neuron's output, Y_i
 - W_{ji} is known as the weight from unit i to unit j
 - $W_{ji} > 0$, synapse is excitatory
 - $W_{ji} < 0$, synapse is inhibitory
 - Note that I_i may be
 - External input
 - The output of some other neuron


Slide 8: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks (FFNN): Overview

A **feedforward neural network** is a biologically inspired **artificial neural network (ANN)**.

It consists of a number of **simple artificial neurons** processing units, arranged in **layers**.

Every unit in a layer is **connected** with **all the units** in the previous layer.




Feedforward Neural Network

Slide 9: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks (FFNN): Overview

- The **artificial neuron units** in a neural network are called **nodes**.
- The links between neuron nodes are not all equal: **Each link** may have a different strength or **weight**.
- These **weights** on these links represent the **knowledge** of a network.
- Data enters at the **inputs** and passes through the **layers** of the network until it **arrives at the outputs**.
- This network is called **feedforward neural networks** because there is **no feedback between layers**.



Feedforward Neural Network

Slide 10: AI Deep Learning: Feedforward Neural Networks (FFNN)


Feedforward Neural Network (FFNN) = Multi-Layer Perceptron (MLP)

Feedforward neural networks are also called

- Deep Feedforward Network

OR

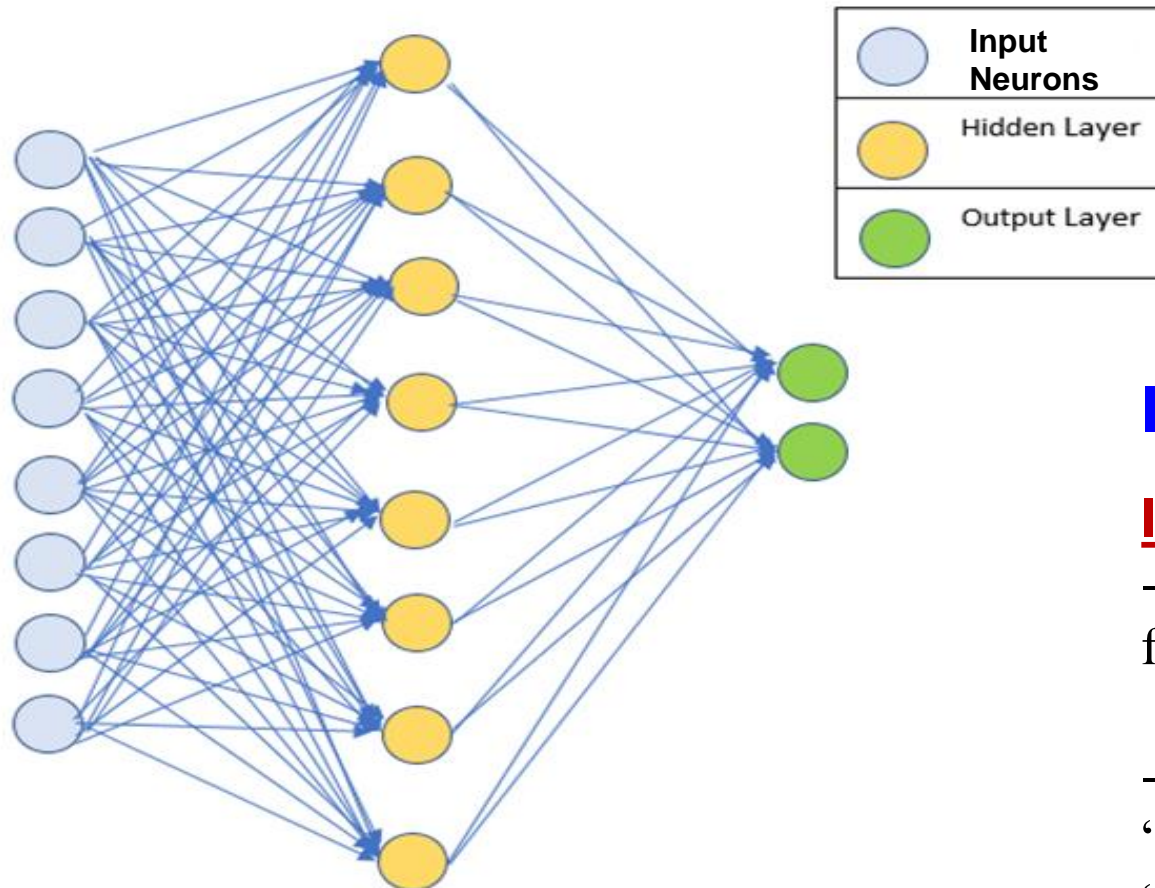
- Multi-Layer Perceptron (MLP)



McCulloch-Pitts neuron model and Feedforward Neural Network (Sources: towardsdatascience.com)

Slide 11: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks (FFNN): Overview



Feedforward Neural Network

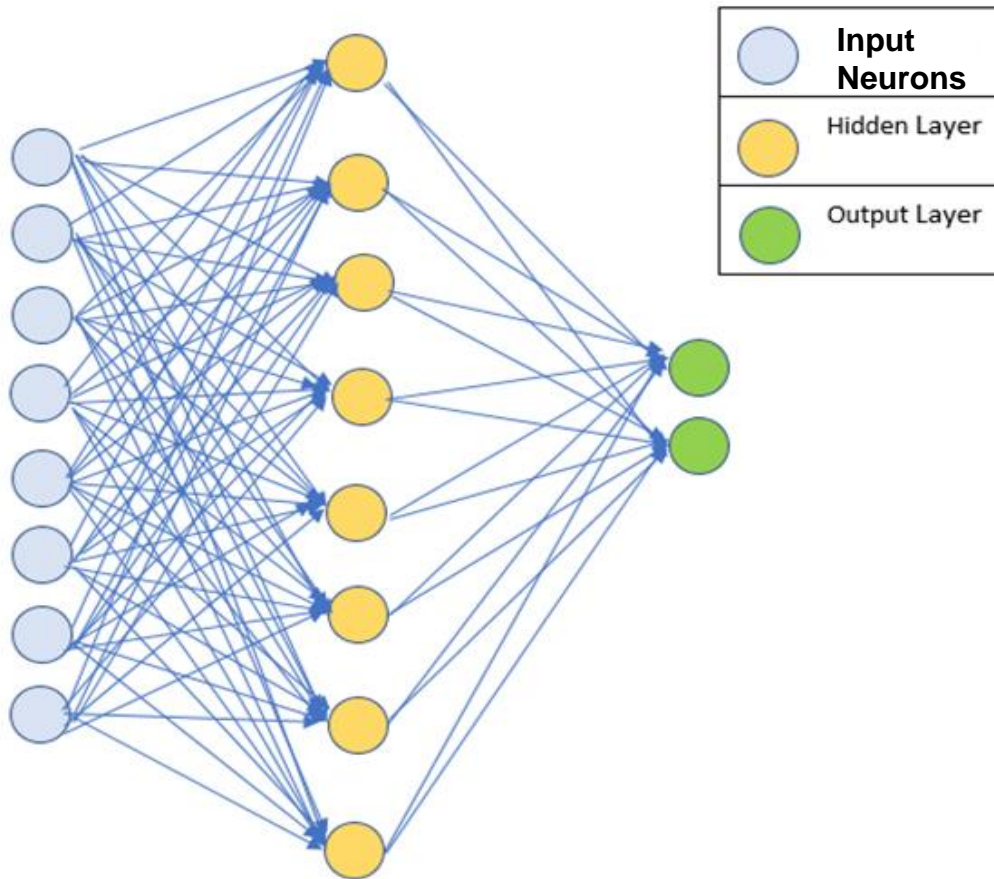
Input Neurons:

IMPORTANT NOTES:

- > Input neurons **do not** belong to any layer of the feedforward neural network.
- > The set of input neurons is sometimes referred to as the “input layer”. However, the **input layer** is viewed as a “**virtual layer**” and **not counted** as a layer of the network.

Slide 12: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks (FFNN): Overview



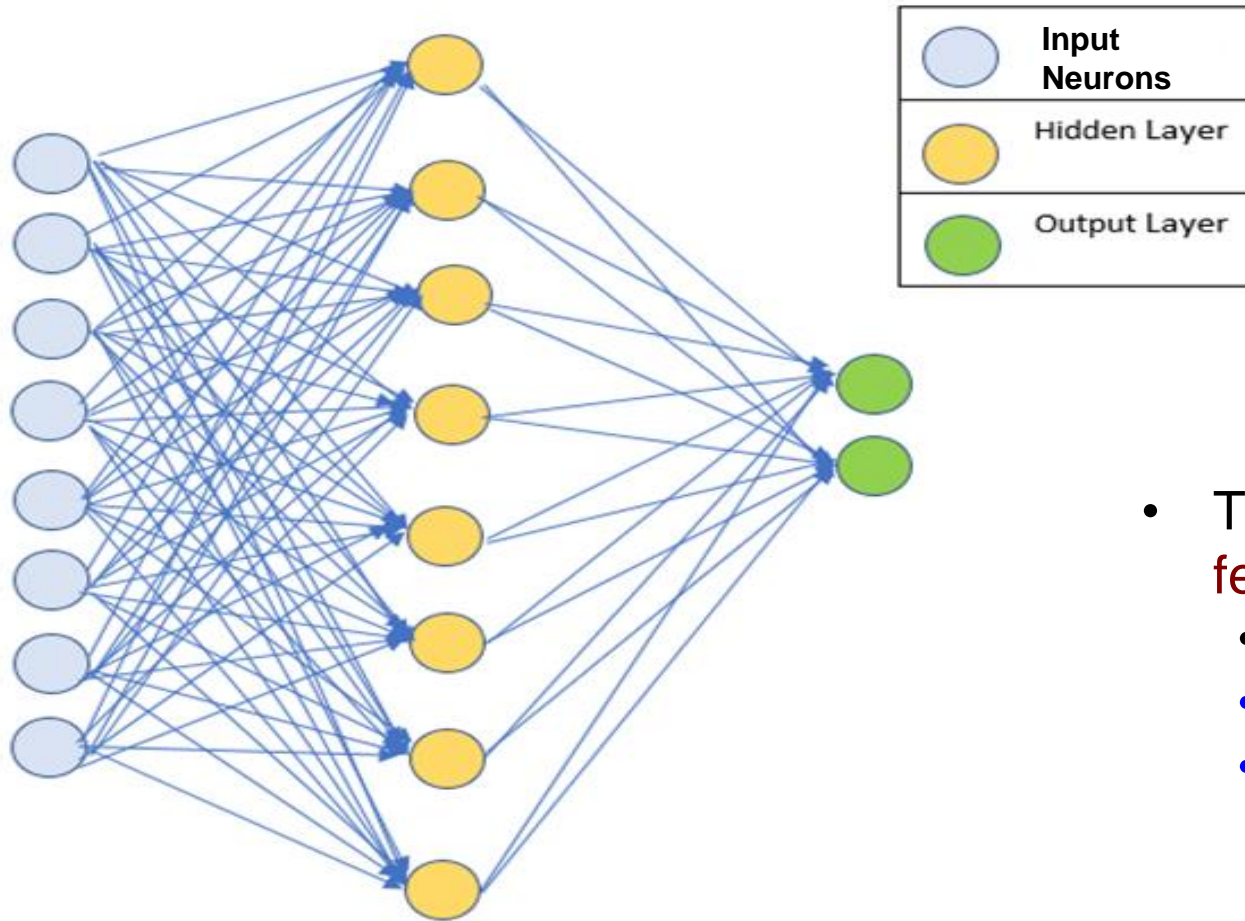
Feedforward Neural Network

Hidden Layers:

- Any layer **between** the **input neurons** and the **output layer** is a **hidden layer**.
- During the training process, the training data **does not show** desired **outputs** of hidden layers.
- A network can consist of **any number** of **hidden layers**.
- Each **hidden layer** can contain **any number** of artificial neuron units/nodes.

Slide 13: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks (FFNN): Overview



- The figure to the left shows an example of a feedforward neural network of two layers:
 - 8 input neurons
 - One **hidden** layer of 8 neurons
 - One **output** layer of 2 neurons


Feedforward Neural Network

Slide 14: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network (FFNN) = Multi-Layer Perceptron (MLP)

In a feedforward neural network:

- Data enters at the input neurons and passes through the layers of the network until it arrives at the output layer, i.e., only one direction, the **forward**.
- There are **no feedback connections** in which **outputs** of the network or **outputs** of each layer are **fed back** into itself.



McCulloch-Pitts neuron model and Feedforward Neural Network (Sources: towardsdatascience.com)

Slide 15: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks:

Feedforward neural networks are the quintessential deep learning models.

- The **goal** of a feedforward network is to **approximate** some function $f^*(x)$.
- For example:
 - For a classifier, $y = f^*(x)$ maps an input x to a category y .
 - A feedforward network defines a mapping $y = f(x; \theta)$ and **learns** the value of the parameter θ that result in the **best function approximation** ([Reference](#)).
- These **networks** are represented by a **composition** of many different functions.
 - For example:
 - We might have three functions $f_1(x)$, $f_2(y)$ where $y = f_1(x)$, and $f_3(z)$ where $z = f_2(y)$.
 - These functions are **connected in a chain**, to form $f(x) = f_3(f_2(f_1(x)))$.
 - In this, $f_1(x)$ is the first layer, $f_2(y)$ is the second layer and $f_3(z)$ is the output layer.

Slide 16: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks: Gradient-Based Learning

- **Gradient Descent** is an **optimization** algorithm .
 - It is based on a **convex** function.
 - It tweaks its parameters **iteratively** to **minimize** a given function to its **local minimum**.
 - It is used in **training** a **machine learning** or **deep learning** model.
- **Gradient descent** is a method to **find** the values of a **cost function's** parameters (coefficients) that **minimize** the function as far as possible.
- What is a **Gradient**?
 - "A gradient measures **how much the output** of a **function changes** if you **change** the **inputs** a **little bit**." - Lex Fridman (MIT)
 - For a machine learning or **deep learning** model, a gradient measures the **change** in all weights with regard to the **change** in cost or error.

Slide 17: AI Deep Learning: Feedforward Neural Networks (FFNN)


Feedforward Neural Networks: Gradient-Based Learning

For the learning process of a machine learning or deep learning model:

- A gradient measures the **change in all weights** with regard to the **change in cost or error**.
 - **Gradient** can be illustrated as the **slope of the cost function**:
 - The **higher** the gradient: The **steeper** the slope → The **faster** a model can **learn**.
 - The **lower** the gradient: The **less steep** the slope → The **slower** a model can **learn**.
 - If the **gradient** is **near zero** or **zero**, the model **stops learning**.
 - The cost function gets to its minimum.
- OR
- This is a **serious well-recorded problem** in training deep learning models.
 - The problem is called **Vanishing Gradient**.

Slide 18: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?



Gradient-Based Learning (Source: Divakar Kapil, medium.com)

Slide 19: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks: Gradient-Based Learning

The **learning process** in feedforward neural networks is **gradient-based**.

However, there are **differences** between **training artificial neural networks** and training a general machine-learning model with **gradient descent**.

For artificial neural networks like feedforward neural networks (FFNN):

- The **nonlinearity** of a neural network causes most **loss functions** to become **non-convex**.
- Neural networks are usually trained by **using iterative, gradient-based optimizers** that **merely drive** the **cost function** to a **very low** value.
 - i.e., **only** as low as possible, **not** the **minimum value** as in the case of convex function.

Slide 20: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks: Optimization & Cost Function

- **Cost function** shows the **difference** between the **approximation** made by the model and the **actual target value**.
- A **cost function** is mostly of form **$C(W, B, S_r, E_r)$** :
 - **W** is the weights of the neural network, **B** is the biases of the network, **S_r** is the input of a single training sample, and **E_r** is the desired output of that training sample.
- Many cost functions can be used for the training process of artificial neural networks:
 - **Mean Squared Error (MSE: a.k.a. Quadratic cost function OR Sum of Squared Errors)**
 - **Cross-entropy cost**
 - Also known as Bernoulli negative log-likelihood and Binary Cross-Entropy
 - **Exponential Cost**
 - **Hellinger distance**
- Among the above mentioned cost functions:
 - **MSE** is very popular for regression problems.
 - **Cross-entropy** is the most often used for classification tasks.

Slide 21: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks: Optimization & Cost Function

- An **Optimizer** or optimization algorithm is used to **minimize the cost function**:
 - It **updates the values of the weights and biases** after **every training** cycle or epoch **until** the cost function reaches the **global optimum**.
- Optimization algorithms (optimizers) are of **two types**:
 - **First Order** Optimization Algorithms
 - **Second Order** Optimization Algorithms
- **First Order** Optimization Algorithms (optimizers):
 - These algorithms **minimize or maximize a cost function** using its **gradient values** with respect to the parameters.
 - The **First Order derivative** tells us **whether the function is decreasing or increasing at a particular point**, in short, it gives the line which is tangent to the surface.

Slide 22: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks: Optimization & Cost Function

- **Second Order** Optimization Algorithms (optimizers):
 - Using **second order derivatives** to **minimize the cost function** and are also called Hessian.
 - Since the **second derivative is costly** to compute, the second order is **not used much**.
 - The second order derivative tells us whether the first derivative is increasing or decreasing which hints at the function's curvature.
- There are many algorithms (optimizers) that can be used for optimization:
 - **Stochastic gradient descent**
 - **Adagrad**
 - **Adam**
 - **RMSProp**.
- Among the above-mentioned optimizers, **Adam** is now the **most popular** optimization algorithm in training artificial neural networks.

Slide 23: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Networks: Optimization & Cost Function

Optimization Algorithms (optimizers): **Adam**

- Is an **extension** to **stochastic gradient descent (SGD)** algorithm.
- Can be used to **update network weights** iterative based on training data.
- Adam was proposed by Diederik Kingma (OpenAI) and Jimmy Ba (University of Toronto)
- The name is derived from “**Adaptive Moment Estimation**,” not an acronym.
- Providing **benefits** on **non-convex optimization** problems
 - **Straightforward to implement**
 - **Computationally efficient**
 - **Little memory requirements**
 - ... MORE
- Combining the **advantages** of **two other popular optimization algorithms**:
 - Adaptive Gradient Algorithm (AdaGrad)
 - Root Mean Square Propagation (RMSProp)

Slide 24: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?

Scenario:

- It is assumed that the **desired output** of the network is y .
- Also assumed that the **neural network** produces an **output y'** .
- The **difference** between the **predicted output** and the **desired output ($y' - y$)** is considered as the **cost or loss** represented by the **cost/loss function**.
 - The loss is high when the neural network makes a lot of mistakes
 - The loss is low when it makes fewer mistakes.

The **goal** of the training process:

- To find the **weights** and **bias** that **minimizes** the **loss function** over the **training set**.

Slide 25: AI Deep Learning: Feedforward Neural Networks (FFNN)

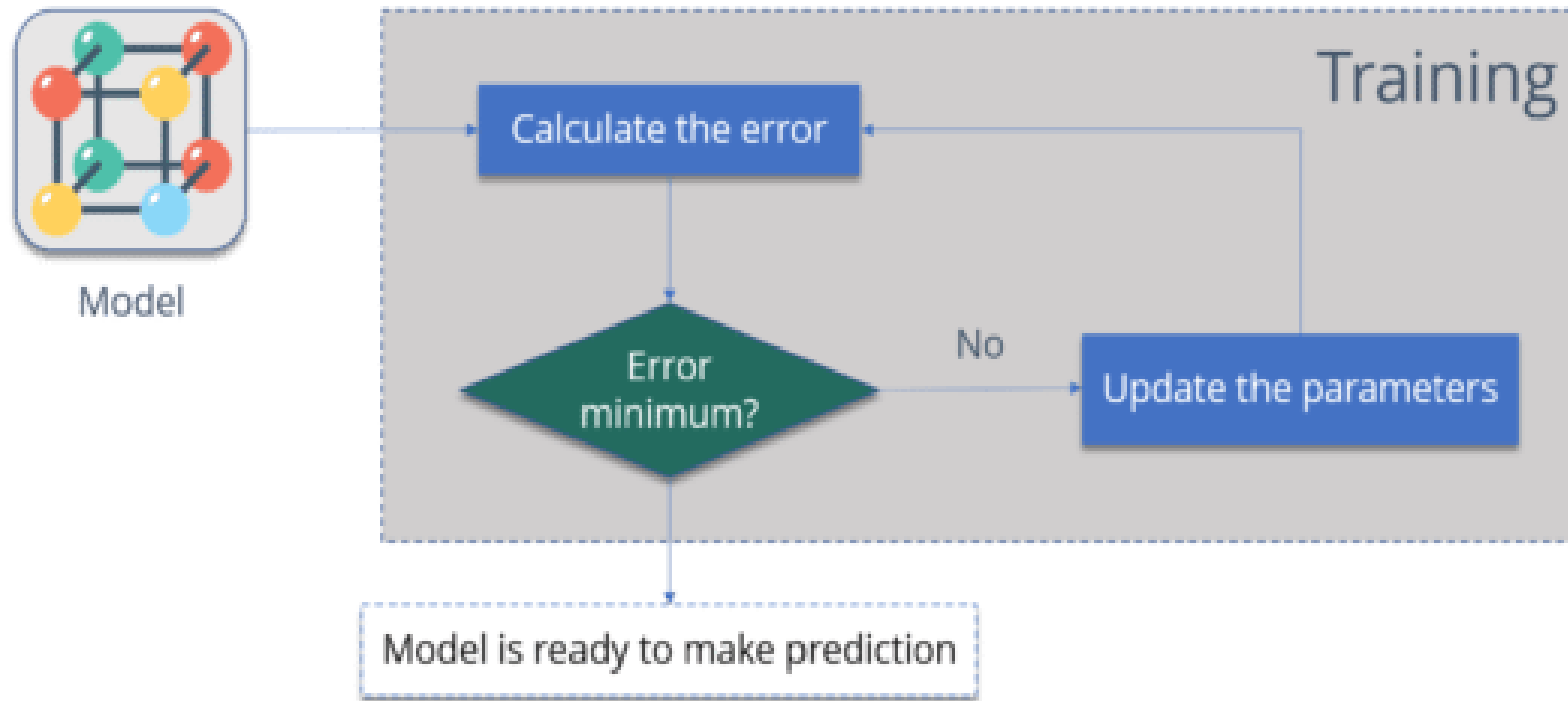
Feedforward Neural Network: How Does It Learn?

Backpropagation Algorithm

- The training is done using the **Backpropagation algorithm**, also called **backprop**.
 - The algorithm **iteratively** passes batches of data through the network and **updating the weights** to **decrease the error**, or cost/loss.
 - The algorithm can do this by **running an optimization algorithm** for deep learning like **Adam** or **Stochastic Gradient Descent (SGD)**.
- The **amount** by which the **weights are changed** is determined by a **parameter** called **Learning Rate**.

Slide 26: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?



Training A Neural Network – Backpropagation (Sources: edureka.co)

Slide 27: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?

Backpropagation Algorithm

- The **backpropagation** algorithm was introduced in the **1970s**.
 - However, its importance was not fully recognized until a famous 1986 paper by **David Rumelhart, Geoffrey Hinton, and Ronald Williams**.
 - That paper describes several neural networks where backpropagation works far faster than earlier approaches to learning.
- Today, the **backpropagation** algorithm is **very popular** as a **technique to train** neural networks.

Slide 28: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?

Backpropagation Algorithm

- **Backpropagation** is the **essence** of **training** neural networks.
 - It is considered the **most effective technique** to train neural networks.
 - It is the **method of fine-tuning the weights** of a neural network based on the error rate obtained in the previous epoch (or iteration).
 - Proper tuning of the weights allows **reducing error rates and making the model reliable** by increasing its generalization.
- **Backpropagation** is a short form for "**backward propagation of errors**."
 - It is a **standard method** of **training** artificial neural networks.
 - This method helps to **calculate the gradient** of a **loss function** with respects to **all the weights** in the network.

Slide 29: AI Deep Learning: Feedforward Neural Networks (FFNN)


Feedforward Neural Network: How Does It Learn?

Backpropagation Algorithm

- Input data is split into training batches.
- First, **initializing** some random value to the weights and **propagating forward**.
- **Forward pass**:
 - The batches are passed from input neurons through the layers of the network to the output layer to produce the outputs.
- **Output comparison**:
 - The outputs are compared with the actual values of the outcomes (or labels) of the data set.
 - The difference, i.e., error, is calculated. **To reduce error**, recursively **propagating backward**.
- **Backward pass**:
 - The difference, i.e., error, cost, or lost, is used to change the weights of the neurons at each layer recursively such that the error decreases gradually.

Slide 30: AI Deep Learning: Feedforward Neural Networks (FFNN)


Feedforward Neural Network: How Does It Learn?



Feedforward Neural Network and Backpropagation Algorithm

Slide 31: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?



Feedforward Neural Network and Backpropagation Algorithm (Sources: guru99.com)

Slide 32: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?

Backpropagation Algorithm

- The **core of backpropagation** algorithm is an expression for the **partial derivative $\partial C / \partial w$** :
 - Is the **partial derivative of the cost function C** with respect to (w.r.t.) **any weight w** (or bias b) in the network.
 - The expression tells us **how quickly** the **cost changes** when we **change** the weights and biases.
- The **partial derivative $\partial C / \partial w$** of the **cost function C** w.r.t. **any weight w** (or bias b)
 - Is similar to the case of **$y = f(t)$**
 - Where **y** is the **distance** that one object is traveling during the time period of t .
 - The **change of distance** per time unit (**speed**): **$y' = f'(t)$** .
 - The **change of the speed** per time unit (**acceleration**): **$y'' = f''(t)$** .

Slide 33: AI Deep Learning: Feedforward Neural Networks (FFNN)

Feedforward Neural Network: How Does It Learn?

Backpropagation Algorithm

- Most prominent **advantages** of **Backpropagation** are:
 - Backpropagation is **fast, simple and easy** to program
 - It has **no parameters to tune** apart from the numbers of input
 - It is a **flexible method** as it does not require prior knowledge about the network
 - It is a **standard method** that generally works well
 - It does **not need any special mention of the features** of the function to be learned.