

Linear Algebra for AI: Deep Learning

Thuan L Nguyen, PhD

Slide 2: Linear Algebra for AI: Deep Learning

1. AI Deep Learning: Linear Algebra and Deep Learning
2. AI Deep Learning: Scalars, Vectors, and Matrices
3. AI Deep Learning: Matrices: Operations: Addition, Subtraction, and Multiplication
4. AI Deep Learning: Matrices and Python Numpy N-Dimensional Arrays
5. AI Deep Learning: Matrices: A Vector of Vectors
6. AI Deep Learning: Matrices: The Concept of Axis
7. AI Deep Learning: Tensors: N-Dimensional Arrays
8. AI Deep Learning: Multi-Class Classifications: One-Hot Coding
9. AI Deep Learning: Activation Functions: ReLU (Rectified Linear Unit) and Softmax

Slide 3: Linear Algebra for AI: Deep Learning

Linear Algebra:

- Linear algebra is a key branch of mathematics.
- An understanding of linear algebra is crucial for deep learning, that is, neural networks.
- Linear Algebra deals with linear systems of equations.
- The focus will be on vectors, matrices, and tensors.
- Using linear algebra, we can describe complicated operations in deep learning.

Slide 4: Linear Algebra for AI: Deep Learning

Scalars, Vectors, and Matrices

- Scalars, vectors, and matrices are the fundamental objects of mathematics:
 - Scalar is represented by a single number or numerical value called magnitude.
 - Vector is an array of numbers assembled in order.
 - A unique index identifies each number.
 - Vector represents a point in space, with each element giving the coordinate along a different axis.
 - Matrices is a two-dimensional array of numbers where each number is identified using two indices (i, j).

Slide 5: Linear Algebra for AI: Deep Learning

Scalars, Vectors, and Matrices: Examples

- Scalars, vectors, and matrices are the fundamental objects of mathematics:
 - Scalar is represented by a single number or numerical value called magnitude.
- For example:
 - 5
 - 16.32
 - -0.128
 - ...


Slide 6: Linear Algebra for AI: Deep Learning

Scalars, Vectors, and Matrices: Examples

- **Scalars, vectors, and matrices** are the fundamental objects of mathematics:
 - **Vector** is an **array of values** assembled in order.
 - A unique **index** identifies each value.
- For example:
 - [1, 3, 5, 7, 9]
 - ['a', 'b', 'c']
 - ["abc", "def"]
- A **vector** can be implemented with a Python Numpy **1D-Array**.



Slide 7: Linear Algebra for AI: Deep Learning



Vector: Index of Elements

- A vector can be implemented with a Python Numpy 1D-Array.
- Usually, the **index** of the **elements** of a vector **starts at 0**.

Slide 8: Linear Algebra for AI: Deep Learning

Scalars, Vectors, and Matrices: Examples

- Scalars, vectors, and matrices are the fundamental objects of mathematics:
 - Matrices** is a two-dimensional array of values where each value is identified using two indices (i, j).
- For example:

$$\begin{bmatrix} 1 & 1 & 1 & -1 \\ 1 & 2 & 4 & 3 \\ 1 & 3 & 9 & 3 \end{bmatrix}$$

$$\begin{bmatrix} 2 & 1 & 3 & 4 & 5 \\ 0 & 0 & 1 & 4 & 2 \\ 4 & 2 & 6 & 8 & 10 \\ 6 & 3 & 14 & 35 & 33 \end{bmatrix}$$

Slide 9: Linear Algebra for AI: Deep Learning

Scalars, Vectors, and Matrices: Examples

- Scalars, vectors, and matrices are the fundamental objects of mathematics:
 - Matrices** is a two-dimensional array of values where each value is identified using two indices (i, j).
- Other examples (Not using square brackets):

$$\begin{pmatrix} 2 & 1 & 2 \\ 1 & -1 & 4 \end{pmatrix}$$

$$B = \begin{pmatrix} 2 & 5 & 14 \\ 1 & 3 & 8 \\ -1 & -2 & -6 \end{pmatrix}$$

Slide 10: Linear Algebra for AI: Deep Learning

Matrices: Operations: Addition & Subtraction

- Two matrices A and B can be added or subtracted if and only if their **dimensions** are the **same** (i.e. both matrices have the same number of rows and columns).
- For example:

$$A = \begin{pmatrix} 1 & 2 & 3 \\ 1 & 0 & 2 \end{pmatrix} \quad \text{and} \quad B = \begin{pmatrix} 2 & 1 & 2 \\ 1 & 0 & 3 \end{pmatrix}$$

$$A + B = \begin{pmatrix} 1 & 2 & 3 \\ 1 & 0 & 2 \end{pmatrix} + \begin{pmatrix} 2 & 1 & 2 \\ 1 & 0 & 3 \end{pmatrix} = \begin{pmatrix} 3 & 3 & 5 \\ 2 & 0 & 5 \end{pmatrix}$$

$$A - B = \begin{pmatrix} 1 & 2 & 3 \\ 1 & 0 & 2 \end{pmatrix} - \begin{pmatrix} 2 & 1 & 2 \\ 1 & 0 & 3 \end{pmatrix} = \begin{pmatrix} -1 & 1 & 1 \\ 0 & 0 & -1 \end{pmatrix}$$

Slide 11: Linear Algebra for AI: Deep Learning

RNN: LSTM: Linear Algebra for Deep Learning:

Vector and Matrix Multiplication: Dot Multiplication

$$(m \times \mathbf{n}) * (\mathbf{n} \times p) \rightarrow (m \times p)$$

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & \end{bmatrix}$$

The "Dot Product" is where we **multiply matching members**, then sum up:

$$(1, 2, 3) \cdot (7, 9, 11) = 1 \times 7 + 2 \times 9 + 3 \times 11 \\ = 58$$

Dot Product of Matrices (Sources: mathsisfun.com)

Slide 12: Linear Algebra for AI: Deep Learning

RNN: LSTM: Linear Algebra for Deep Learning:

Vector and Matrix Multiplication: Dot Multiplication

$$(m \times \mathbf{n}) * (\mathbf{n} \times p) \rightarrow (m \times p)$$

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & 64 \\ 64 & 74 \end{bmatrix}$$

$$(1, 2, 3) \cdot (8, 10, 12) = 1 \times 8 + 2 \times 10 + 3 \times 12 \\ = 64$$

Dot Product of Matrices (Sources: mathsisfun.com)

Slide 13: Linear Algebra for AI: Deep Learning

RNN: LSTM: Linear Algebra for Deep Learning:

Vector and Matrix Multiplication: Dot Multiplication

We can do the same thing for the **2nd row** and **1st column**:

$$(4, 5, 6) \bullet (7, 9, 11) = 4 \times 7 + 5 \times 9 + 6 \times 11 \\ = 139$$

And for the **2nd row** and **2nd column**:

$$(4, 5, 6) \bullet (8, 10, 12) = 4 \times 8 + 5 \times 10 + 6 \times 12 \\ = 154$$

$$(m \times \mathbf{n}) * (\mathbf{n} \times p) \rightarrow (m \times p)$$

And we get:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & 64 \\ 139 & 154 \end{bmatrix} \checkmark$$

Dot Product of Matrices (Sources: mathsisfun.com)

Slide 14: Linear Algebra for AI: Deep Learning

RNN: LSTM: Linear Algebra for Deep Learning:

Vector and Matrix Multiplication: Dot Multiplication

$$(m \times \mathbf{n}) * (\mathbf{n} \times p) \rightarrow (m \times p)$$

$$A = \begin{bmatrix} 1 & 3 \\ -1 & 0 \end{bmatrix} B = \begin{bmatrix} 2 & 1 & 1 \\ -1 & 2 & 4 \end{bmatrix}$$


$$(2 \times \mathbf{2}) * (\mathbf{2} \times 3) \rightarrow (2 \times 3)$$

$$AB = \begin{bmatrix} 1 \cdot 2 + 3 \cdot -1 & 1 \cdot 1 + 3 \cdot 2 & 1 \cdot 1 + 3 \cdot 4 \\ -1 \cdot 2 + 0 \cdot -1 & -1 \cdot 1 + 0 \cdot 2 & -1 \cdot 1 + 0 \cdot 4 \end{bmatrix} = \begin{bmatrix} -1 & 7 & 13 \\ -2 & -1 & -1 \end{bmatrix}$$

Slide 15: Linear Algebra for AI: Deep Learning

Matrices & Python Numpy N-Dimensional Arrays

- A **matrix** can be viewed as a **vector of vectors**; each **vector** is one **row**.



Slide 16: Linear Algebra for AI: Deep Learning

Matrices & Python Numpy N-Dimensional Arrays

- A **matrix** can be viewed as a **vector of vectors**; each **vector** is one **row**.

$$\begin{bmatrix} 1 & 1 & 1 & -1 \\ 1 & 2 & 4 & 3 \\ 1 & 3 & 9 & 3 \end{bmatrix}$$

- The matrix (a vector of vectors) has three vector elements:
 - Index 0: $[1, 1, 1, -1]$
 - Index 1: $[1, 2, 4, 3]$
 - Index 2: $[1, 3, 9, 3]$
- The matrix can be displayed in another format:
 $[[1, 1, 1, -1], [1, 2, 4, 3], [1, 3, 9, 3]]$

Slide 17: Linear Algebra for AI: Deep Learning

Matrices & Python Numpy N-Dimensional Arrays: The Concept of Axis

IMPORTANT NOTES: Axis = 0

- Matrix: A vector of vectors

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15

- **Axis 0**:
 - First axis
 - Direction: Moving from one vector to another vector
 - For example: Element-wise along Axis 0
 - $1 \rightarrow 6 \rightarrow 11$
 - $2 \rightarrow 7 \rightarrow 12$
 - ...
 - $5 \rightarrow 10 \rightarrow 15$

Slide 18: Linear Algebra for AI: Deep Learning

Matrices & Python Numpy N-Dimensional Arrays: The Concept of Axis

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15

- **Axis 1:**
 - Second axis
 - Direction: Moving from one element to another element in the same vector
 - For example: Element-wise along Axis 1
 - $1 \rightarrow 2 \rightarrow 3 \rightarrow \dots$
 - $6 \rightarrow 7 \rightarrow 8 \rightarrow \dots$
 - \dots
 - $11 \rightarrow 12 \rightarrow \dots$

Slide 19: Linear Algebra for AI: Deep Learning

Matrices & Python Numpy N-Dimensional Arrays

- A **matrix** can be viewed as a **vector of vectors**.
- Instead of displaying a **matrix** in the form of **rows x columns**, it is possible to show the **matrix** as a **vector** of which **each element** is one **row** of the **matrix**.

$$A = \begin{bmatrix} 1 & 3 \\ -1 & 0 \end{bmatrix} = \text{[[1, 3], [-1, 0]]}$$

$$B = \begin{bmatrix} 2 & 1 & 1 \\ -1 & 2 & 4 \end{bmatrix} = \text{[[2, 1, 1], [-1, 2, 4]]}$$

Slide 20: Linear Algebra for AI: Deep Learning

Flattening a Numpy N-Dimensional Array into a Vector (an 1D-Array)

IMPORTANT NOTES:

- $[[1, 3], [-1, 0]] \neq [1, 3, -1, 0]$
 - $[1, 3, -1, 0]$ is a **vector**, i.e., a 1D-array, that has 4 elements; each element is a scalar
 - $[[1, 3], [-1, 0]]$ is a **matrix**, i.e., a 2D-array. That has 2 elements; each element is a vector.

To **flatten** a Numpy N-Dimensional array, say **ND-array**, is to **collapse** all the **elements** of ND-array into a vector, or a one-dimensional array, **1D-array**.

Slide 21: Linear Algebra for AI: Deep Learning

Flattening a Numpy N-Dimensional Array into a Vector (an 1D-Array)


For example:

- $\begin{bmatrix} [1, 3], [-1, 0] \end{bmatrix}$ ----- flattening \rightarrow $[1, 3, -1, 0]$
- $\begin{bmatrix} [2, 1, 1], [-1, 2, 4], [1, 3, 1] \end{bmatrix}$ ----- flattening \rightarrow $[2, 1, 1, -1, 2, 4, 1, 3, 1]$

Slide 22: Linear Algebra for AI: Deep Learning

Tensors: N-Dimensional Arrays

- In mathematics, a tensor is a geometric object:
 - This object can map geometric vectors, scalars, and other tensors to a resulting tensor in a multi-linear manner.
- Geometric vectors and scalars are considered as the simplest tensors.
- The dimension of a tensor is called its rank.
- The image to the left describes a tensor of rank 2, i.e., 2nd-order tensor, in a 3D Cartesian coordinate system:



Tensor (Source: Wikipedia)

Slide 23: Linear Algebra for AI: Deep Learning

Tensors: N-Dimensional Arrays

- To be simple, a tensor is often thought of as a generalized matrix. It could be:
 - A 0-D matrix: a single number, i.e., a scalar
 - A 1-D matrix: a vector
 - A 2-D matrix
 - A higher dimensional array, i.e., a N-Dimensional array
- A tensor can be implemented as a Python Numpy N-Dimensional array.

Slide 24: Linear Algebra for AI: Deep Learning

Tensors: N-Dimensional Arrays

- A scalar is a tensor ($f : \mathbb{R} \rightarrow \mathbb{R}, f(e_1) = c$)
- A vector is a tensor ($f : \mathbb{R}^n \rightarrow \mathbb{R}, f(e_i) = v_i$)
- A matrix is a tensor ($f : \mathbb{R}^n \times \mathbb{R}^m \rightarrow \mathbb{R}, f(e_i, e_j) = A_{ij}$)
- Common to have fixed basis, **so a tensor can be represented as a multidimensional array of numbers.**

Slide 25: Linear Algebra for AI: Deep Learning

Multi-Class Classification: Encoding Classes with One-Hot Coding

- One of the major problems with Machine Learning:
 - It is a daunting task to work directly with categorical data.
- Computers are powerful and efficient when dealing with numbers.
 - So, it would be much better to convert our input data (if it is not in numeric formats) to numbers.

Slide 26: Linear Algebra for AI: Deep Learning

Multi-Class Classification: Encoding Classes with One-Hot Coding

- What is category data?
 - Categorical data are variables that contain label values rather than numeric values.
 - The number of possible values is often limited to a fixed set.
 - Categorical variables are often called nominal.
- Some examples include:
 - A “pet” variable with the values: “dog” and “cat”.
 - A “color” variable with the values: “red”, “green” and “blue”.
 - A “place” variable with the values: “first”, “second” and “third”.
- Each value represents a different category.
- Some categories may have a natural relationship to each other, such as a natural ordering.
 - The “place” variable above does have a natural ordering of values.
 - This type of categorical variable is called an ordinal variable.

Slide 27: Linear Algebra for AI: Deep Learning

Multi-Class Classification: Encoding Classes with One-Hot Coding

What is the problem with category data?

- Some algorithms can work with categorical data directly.
 - For example, a decision tree can be learned directly from categorical data with no data transform required (this depends on the specific implementation).
- Many machine learning algorithms cannot operate on label data directly.
 - They require all input variables and output variables to be numeric.
 - In general, this is mostly a constraint of the efficient implementation of machine learning algorithms rather than hard limitations on the algorithms themselves.
- This means that categorical data must be converted to a numerical form.
 - If the categorical variable is an output variable, it is desirable to convert predictions by the model back into a categorical form in order to present them or use them in some application.

Slide 28: Linear Algebra for AI: Deep Learning

Multi-Class Classification: Encoding Classes with One-Hot Coding

How to Convert Categorical Data to Numerical Data?

- The **conversion** is done in **two steps**:
 - Integer coding
 - One-hot coding

First Step: Integer Coding

- First, each unique category value is assigned an integer value.
 - For example:
 - “Yes” = 1, “No” = 0
 - “*red*” = 0, “*green*” = 1, “*blue*” = 2.
 - “sedan” = 1, “truck” = 2; “sport-utility” = 3.
- This is called a label encoding or an integer encoding and is easily reversible.
- The integers have a natural ordered relationship that is good for ordinal and boolean variables
 - In this case, integer coding is **good enough**, e.g., “Yes” = 1; “No” = 0.

Slide 29: Linear Algebra for AI: Deep Learning

Multi-Class Classification: Encoding Classes with One-Hot Coding

How to Convert Categorical Data to Numerical Data?

Second Step: One-Hot Coding

- For **categorical variables** where **no** such ordinal relationship exists:
 - The **integer encoding** is **not** enough.
- In fact, using the **integer encoding** and allowing the model to assume a natural ordering between categories may result in **poor performance** or **unexpected results**.
 - Like predictions halfway between categories
- In this case, a **one-hot encoding** can be applied to the **integer representation**.
 - A **vector (1D array) of binary values** is used to represent each unique integer value.
 - For example, for three colors, there are 3 categories that needs 3 vectors of binary values:
 - “Red” = 0: $[1, 0, 0] \rightarrow$ the value at the **index 0** = 1; all the values at other indices = 0
 - “Green” = 1: $[0, 1, 0] \rightarrow$ the value at the **index 1** = 1; all the values at other indices = 0
 - “Blue” = 2: $[0, 0, 1] \rightarrow$ the value at the **index 2** = 1; all the values at other indices = 0

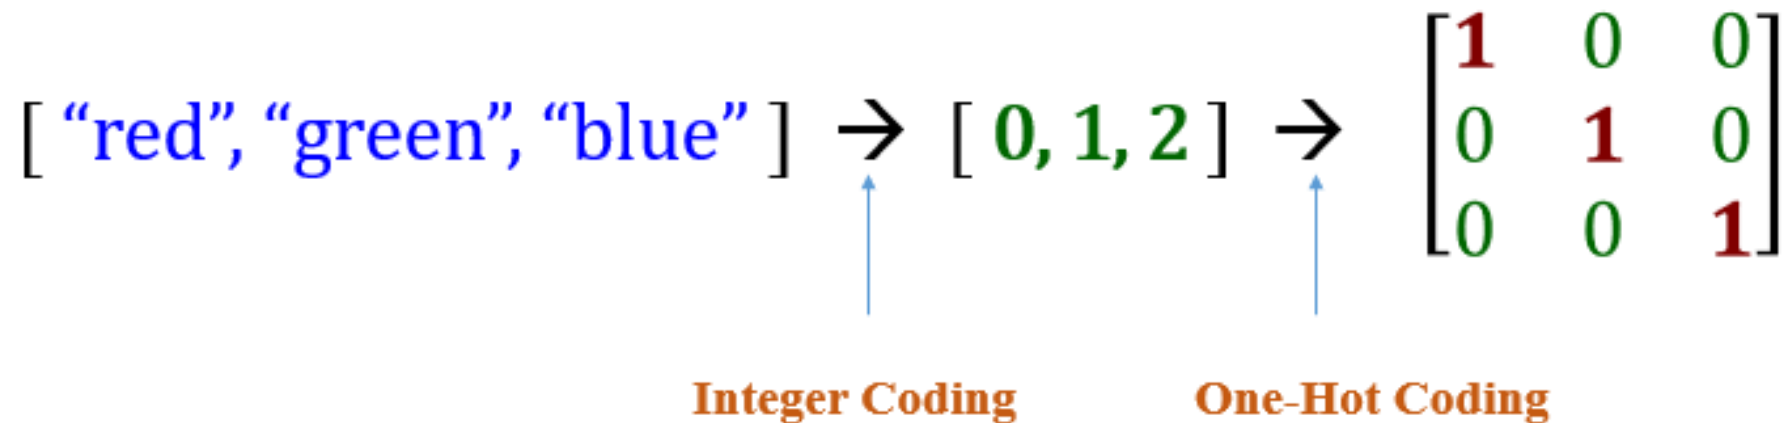
Slide 30: Linear Algebra for AI: Deep Learning

Multi-Class Classification: Encoding Classes with One-Hot Coding

How to Convert Categorical Data to Numerical Data?

Second Step: **One-Hot Coding**

- With the one-hot coding, the set of categorical values (“red”, “green”, “blue”) is transformed from a **vector** (1D array) into a **matrix** (2D array = vector of vectors):



Slide 31: Linear Algebra for AI: Deep Learning

Multi-Class Classification: Encoding Classes with One-Hot Coding

How to Convert Categorical Data to Numerical Data?

Second Step: One-Hot Coding

- With the one-hot coding, the set of categorical values (“red”, “green”, “blue”) is transformed from a **vector** (1D array) into a **matrix** (2D array = vector of vectors).

IMPORTANT NOTES:


--> If the class attribute of a dataset has *only two categorical values*, e.g., “**Yes**” and “**No**”, “Approved” and “Rejected”, etc., in most cases, the *integer coding is enough*.

--> If the class attribute of a dataset has *more than two categorical values*, *even if that they are already in numeric formats like integers*, it should be better to *encode them using one-hot coding* to get better performance in AI machine learning and deep learning.

Slide 32: Linear Algebra for AI: Deep Learning

AI: Deep Learning: Activation Functions

- Many activation functions can be used in the layers of a neural networks.
 - Two most popular ones: **ReLU** and **Softmax**
- ReLU (Rectified Linear Unit) function:
 - The Rectified Linear Unit (ReLU) has become very popular in the last few years.
 - It computes the function $f(x)=\max(0,x)$.
 - In other words, the activation is simply thresholded at zero (see image to the right).



Activation Function ReLU (Source: Stanford)

Slide 33: Linear Algebra for AI: Deep Learning

AI: Deep Learning: Activation Functions

- **Softmax** function:
 - In mathematics, the **softmax** function, also known as **softargmax** or **normalized exponential function** is a function that takes as **input** a vector of K real numbers, and **normalizes** it into a **probability distribution consisting of K probabilities**. (Source: Wikipedia)
 - It means that prior to applying softmax, some vector components could be negative, or greater than one; and might not sum to 1; but after applying softmax, each component will be in the interval $(0,1)$, and the components will add up to 1, so that they can be interpreted as probabilities.
 - Furthermore, the larger input components will correspond to larger probabilities. Softmax is often used in neural networks, to map the non-normalized output of a network to a probability distribution over predicted output classes.

The standard (unit) softmax function $\sigma : \mathbb{R}^K \rightarrow \mathbb{R}^K$ is defined by the formula

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K \text{ and } \mathbf{z} = (z_1, \dots, z_K) \in \mathbb{R}^K$$

- The **softmax** function is often used in the **final layer of a neural network-based classifier**.
 - Such networks are commonly trained under a log loss (or cross-entropy) regime, giving a non-linear variant of multinomial logistic regression. (Source: Wikipedia)