

# AI Deep Learning: Introduction to TensorFlow

Thuan L Nguyen, PhD

# AI Deep Learning: Introduction to TensorFlow

1. TensorFlow: What is It? & Why Now?
2. TensorFlow: What are Tensors?
3. TensorFlow Programs: The Basics
4. TensorFlow Programs: `tf.constant` and `tf.Variable`
5. TensorFlow Programs: `tf.placeholder` & Feed Dictionaries
6. TensorFlow Programs: Graph and Session
7. TensorFlow & Python Numpy Library
8. TensorFlow: Artificial Neural Networks with TensorFlow

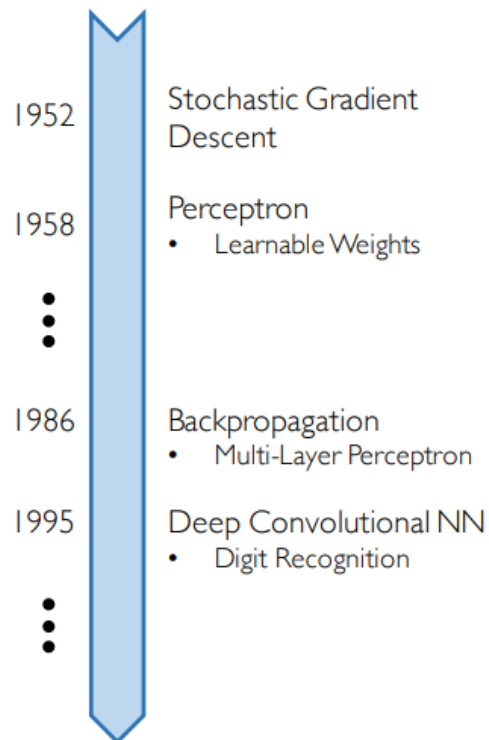
# AI Deep Learning: Introduction to TensorFlow



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: What is It?

- TensorFlow is a very popular library for deep learning computation.
- It is open-sourced by Google.



Sources: MIT

Neural Networks date back decades, so why the resurgence?

### 1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



### 2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



### 3. Software

- Improved Techniques
- New Models
- Toolboxes



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Why?

- Are there any other frameworks that can be used for deep learning computation?
- Yes. Actually, too many!



**Denny Britz** @dennybritz · 25 Dec 2017



I'm going through my newsletters to write up a year-end summary of developments and achievements in AI.



Fun fact: Almost every week, a company released a new generic or task-specific Deep Learning “framework” 😄

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Why?

- Are there any other frameworks that can be used for deep learning computation?
- Yes. Actually, too many!

- Torch
- Caffe
- Theano (Keras, Lasagne)
- CuDNN
- Tensorflow
- Mxnet
- Etc.



# AI Deep Learning: Introduction to TensorFlow


## TensorFlow: Why?

- TensorFlow is a suite of software that has powerful features for deep learning.
  - The features can be used in both academic research and industry production.
- Some major benefits of using TensorFlow:
  - It works with all the cool languages. TensorFlow works with Python, C++, Java, R, and Go.
  - TensorFlow works on multiple platforms, even mobile and distributed.
  - It is supported by all cloud providers – AWS, Google, and Azure.
  - Keras, a high-level neural network API, has been integrated with TensorFlow.
  - It has better computational graph visualizations because it is native.
  - TensorFlow allows model deployment and ease of use in production.
  - TensorFlow has very good community support.
  - TensorFlow is more than a software library; it is a suite of software that includes TensorFlow, TensorBoard, and TensorServing.

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Why?

- TensorFlow is widely used for deep learning projects.
  - It is currently the top deep learning framework in research and production.




Sources: Jeff Hale – Towards Data Science



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Why?

- TensorFlow is widely used for deep learning projects.
  - It is currently the top deep learning framework in research and production.




Sources: Andrej Karpathy and Kiwisoft

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Why?

- TensorFlow is widely used for deep learning projects.
  - It is currently the top deep learning framework in research and production.



Sources: Andrej Karpathy and Kiwisoft

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: What Is a Tensor?

- TensorFlow provides primitives for defining functions on **tensors** and automatically computing their derivatives.
- A scalar is a tensor ( $f : \mathbb{R} \rightarrow \mathbb{R}, f(e_1) = c$ )
- A vector is a tensor ( $f : \mathbb{R}^n \rightarrow \mathbb{R}, f(e_i) = v_i$ )
- A matrix is a tensor ( $f : \mathbb{R}^n \times \mathbb{R}^m \rightarrow \mathbb{R}, f(e_i, e_j) = A_{ij}$ )
- Common to have fixed basis, **so a tensor can be represented as a multidimensional array of numbers.**

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Constants and Variables

- In TensorFlow programming: Three basic elements are constants, variables, and placeholder.
- Constants (`tf.constant`):
  - a constant has a constant value and once you set it, it cannot be changed.
    - `a = tf.constant(2, tf.int16)`
    - `b = tf.constant(4, tf.float32)`
    - `c = tf.constant(8, tf.float32)`
- Variables (`tf.Variable`):
  - The value of a variable can be changed after it has been set, but its type and shape cannot be changed.
    - `d = tf.Variable(2, tf.int16)`
    - `e = tf.Variable(4, tf.float32)`
    - `f = tf.Variable(8, tf.float32)`

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Placeholders and Feed\_Dicts

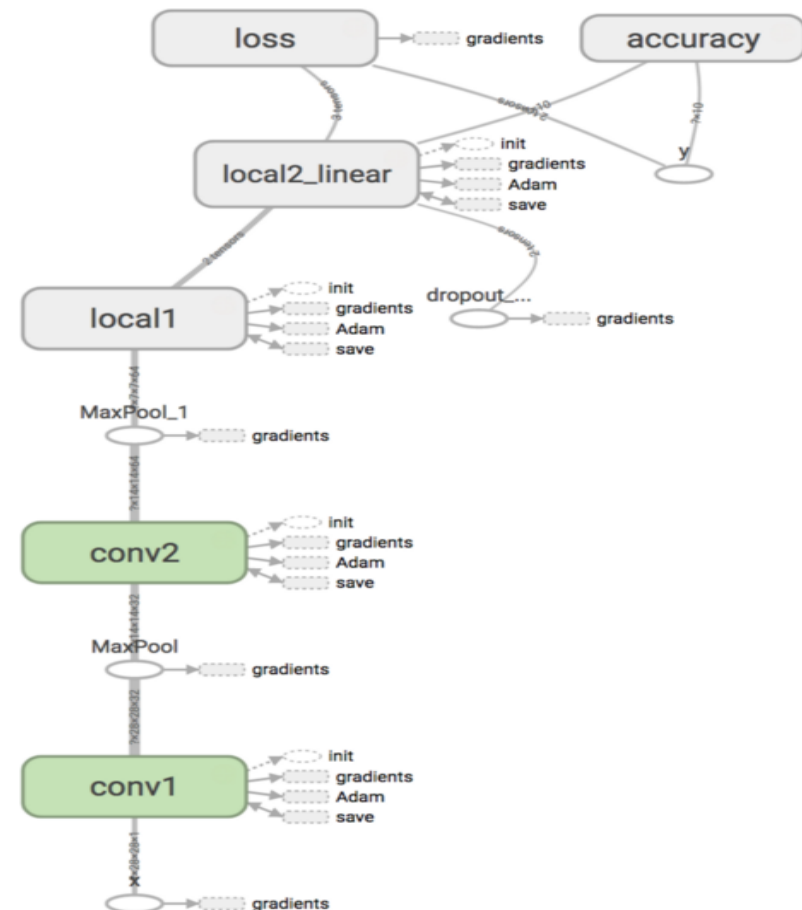
- Tensorflow also has **placeholders**:
  - Placeholders **do not require an initial value**.
  - Placeholders only serve to **allocate the necessary amount of memory**.
    - During a TensorFlow **session**, these placeholder can be filled in with (external) data using a `feed_dict`.
      - *Feed-dict: a python dictionary used to feed data into the computation process*
- Placeholders:
  - Placeholders only serve to **allocate the necessary amount of memory**.
    - `point1 = tf.placeholder(tf.float32, shape=(1, 2))`
    - `point2 = tf.placeholder(tf.float32, shape=(1, 2))`

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Graphs and Session

- In Tensorflow, **all the variables and the operations** done on these variables are **saved in a graph**:
  - After a graph has been built, it **contains all of the computational steps** necessary for the model.
  - The graph can be run (executed) within a **Tensorflow session**.
  - The session distributes all of the computations across the available CPU and GPU resources.
- **TensorFlow programs** are usually **structured into two phases**:
  - A **construction phase**, that **assembles a graph**.
  - And an **execution phase** that uses a session to **execute operations in the graph**.
  - All computations add **nodes** to **global default graph**.


Main Graph



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Graph

- Graphs are sets of connected nodes (vertices).
  - The connections are referred to as edges.
  - In TensorFlow: Each node is an operation with possible inputs that can supply some output.



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Graphs and Session

- Graphs are sets of connected nodes (vertices).
  - The connections are referred to as edges.
  - In TensorFlow: Each node is an operation with possible inputs that can supply some output.
- Examples of graphs and the session:

```
graph = tf.Graph()
with graph.as_default():
    a = tf.Variable(8, tf.float32)
    b = tf.Variable(tf.zeros([2,2], tf.float32))

with tf.Session(graph=graph) as sess:
    anInitializer = tf.global_variables_initializer()
    sess.run(anInitializer)
    print(sess.run(a))
    print(sess.run(b))
```



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Session

- A **Session** object encapsulates the environment in which Tensor objects are evaluated

```
In [20]: a = tf.constant(5.0)
```

```
In [21]: b = tf.constant(6.0)
```

```
In [22]: c = a * b
```

```
In [23]: with tf.Session() as sess:
```

```
.....:     print(sess.run(c))
```

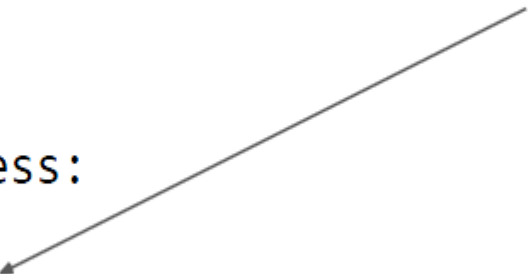
```
.....:     print(c.eval())
```

```
.....:
```

```
30.0
```

```
30.0
```

*c.eval() is just syntactic sugar for  
sess.run(c) in the currently active  
session!*



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Variables


- When you train a model you use variables to hold and update parameters.
  - Variables are in-memory buffers containing tensors

```
In [32]: W1 = tf.ones((2,2))
```

```
In [33]: W2 = tf.Variable(tf.zeros((2,2)), name="weights")
```

```
In [34]: with tf.Session() as sess:
          print(sess.run(W1))
          sess.run(tf.initialize_all_variables())
          print(sess.run(W2))
```

```
.....:
[[ 1.  1.]
 [ 1.  1.]]
[[ 0.  0.]
 [ 0.  0.]]
```



*Note the initialization step `tf.initialize_all_variables()`*


# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Variables

- TensorFlow **variables** must be **initialized** before they can be used.
  - Contrast with constant tensors.

```
In [38]: W = tf.Variable(tf.zeros((2,2)), name="weights")
```


*Variable objects can be  
initialized from constants or  
random values*



```
In [39]: R = tf.Variable(tf.random_normal((2,2)), name="random_weights")
```

```
In [40]: with tf.Session() as sess:  
.....:     sess.run(tf.initialize_all_variables())  
.....:     print(sess.run(W))  
.....:     print(sess.run(R))  
.....:
```

*Initializes all variables with  
specified values.*



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Variables

- TensorFlow variables can be updated via its state

```
In [63]: state = tf.Variable(0, name="counter")
```

```
In [64]: new_value = tf.add(state, tf.constant(1))
```

*Roughly new\_value = state + 1*

```
In [65]: update = tf.assign(state, new_value)
```

*Roughly state = new\_value*

```
In [66]: with tf.Session() as sess:
.....:     sess.run(tf.initialize_all_variables())
.....:     print(sess.run(state))
.....:     for _ in range(3):
.....:         sess.run(update)
.....:         print(sess.run(state))
.....:
```

*Roughly*

```
state = 0
print(state)
for _ in range(3):
    state = state + 1
    print(state)
```

0  
1  
2  
3

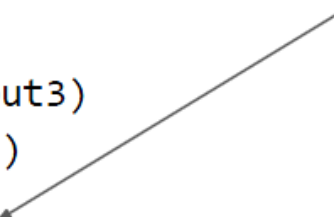
# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Variables

- The states of TensorFlow variables can be retrieved by running a session.

```
In [82]: input1 = tf.constant(3.0)
In [83]: input2 = tf.constant(2.0)
In [84]: input3 = tf.constant(5.0)
In [85]: intermed = tf.add(input2, input3)
In [86]: mul = tf.mul(input1, intermed)
In [87]: with tf.Session() as sess:
.....:     result = sess.run([mul, intermed])
.....:     print(result)
.....:
[21.0, 7.0]
```

Calling `sess.run(var)` on a `tf.Session()` object retrieves its value. Can retrieve multiple variables simultaneously with `sess.run([var1, var2])` (See *Fetches* in TF docs)



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Inputting Data Directly from Numpy

- It is possible to feed data into a TensorFlow program by converting data from Numpy.

```
In [93]: a = np.zeros((3,3))
In [94]: ta = tf.convert_to_tensor(a)
In [95]: with tf.Session() as sess:
        ....:     print(sess.run(ta))
        ....:
[[ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
```

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Inputting Data Using Placeholders & Feed Dictionaries

- It is possible to feed data into a TensorFlow program using `tf.placeholder` and feed dictionaries.

```
In [96]: input1 = tf.placeholder(tf.float32)
```

```
In [97]: input2 = tf.placeholder(tf.float32)
```

```
In [98]: output = tf.mul(input1, input2)
```

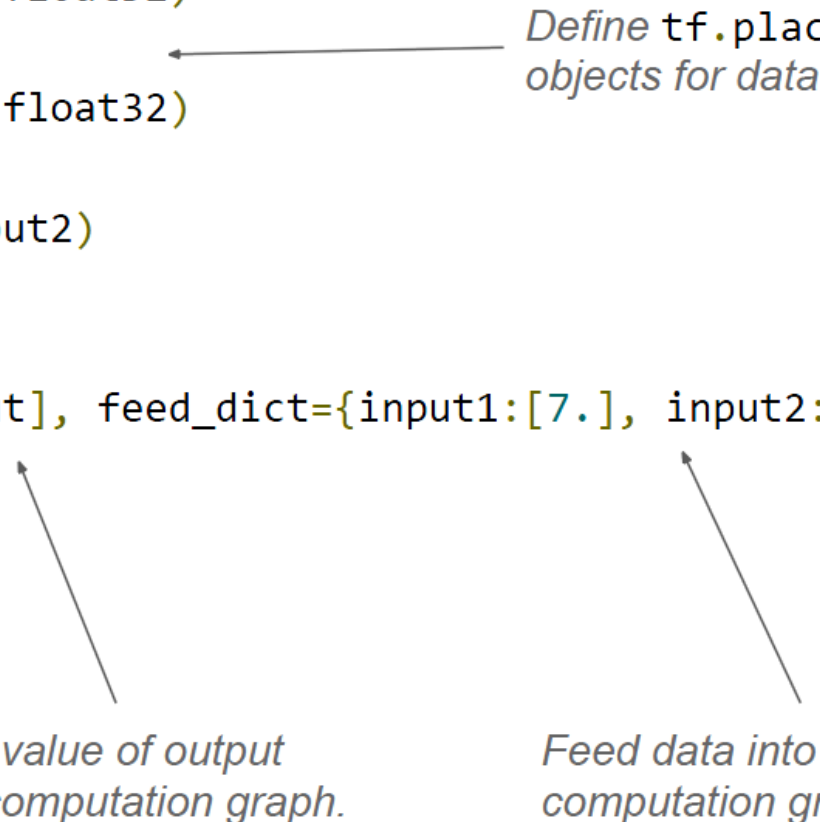
```
In [99]: with tf.Session() as sess:
```

```
.....:     print(sess.run([output], feed_dict={input1:[7.], input2:[2.]}))
```

```
.....:
```

```
[array([ 14.], dtype=float32)]
```

*Define `tf.placeholder`  
objects for data entry.*



The diagram consists of three arrows. One arrow points from the text 'Define tf.placeholder objects for data entry.' to the `tf.placeholder` calls in lines [96] and [97]. A second arrow points from the `feed_dict` parameter in the `sess.run` call to the dictionary `{input1:[7.], input2:[2.]}`. A third arrow points from the `output` variable in the `sess.run` call to the `[output]` list in the same call.


*Fetch value of output  
from computation graph.*

*Feed data into  
computation graph.*

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Inputting Data Using Placeholders & Feed Dictionaries

- It is possible to feed data into a TensorFlow program using `tf.placeholder` and feed dictionaries.
  - Placeholders are initially empty.
  - They are used to feed in the actual training examples.
  - However they do need a declared expected data type (`tf.float32`) with an optional shape argument.





# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: vs Numpy

- TensorFlow seems to be similar to Python Numpy library.
  - Are they the same?
- Few people make this comparison, but TensorFlow and Numpy are quite similar. (Both are N-d array libraries!)
  - Numpy has Ndarray support, but doesn't offer methods to create tensor functions and automatically compute derivatives (+ no GPU support).



VS



# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: vs Numpy

- TensorFlow seems to be similar to Python Numpy library.
- Are they the same?

Numpy	TensorFlow
<code>a = np.zeros((2,2)); b = np.ones((2,2))</code>	<code>a = tf.zeros((2,2)), b = tf.ones((2,2))</code>
<code>np.sum(b, axis=1)</code>	<code>tf.reduce_sum(a, reduction_indices=[1])</code>
<code>a.shape</code>	<code>a.get_shape()</code>
<code>np.reshape(a, (1,4))</code>	<code>tf.reshape(a, (1,4))</code>
<code>b * 5 + 1</code>	<code>b * 5 + 1</code>
<code>np.dot(a,b)</code>	<code>tf.matmul(a, b)</code>
<code>a[0,0], a[:,0], a[0,:]</code>	<code>a[0,0], a[:,0], a[0,:]</code>

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: vs Numpy

- TensorFlow seems to be similar to Python Numpy library.
- Are they the same?

```
In [37]: a = np.zeros((2,2))
```

```
In [38]: ta = tf.zeros((2,2))
```

```
In [39]: print(a)
```

```
[[ 0.  0.]  
 [ 0.  0.]]
```

```
In [40]: print(ta)
```

```
Tensor("zeros_1:0", shape=(2, 2), dtype=float32)
```

```
In [41]: print(ta.eval())
```

```
[[ 0.  0.]  
 [ 0.  0.]]
```

*TensorFlow computations define a **computation graph** that has no numerical value until evaluated!*

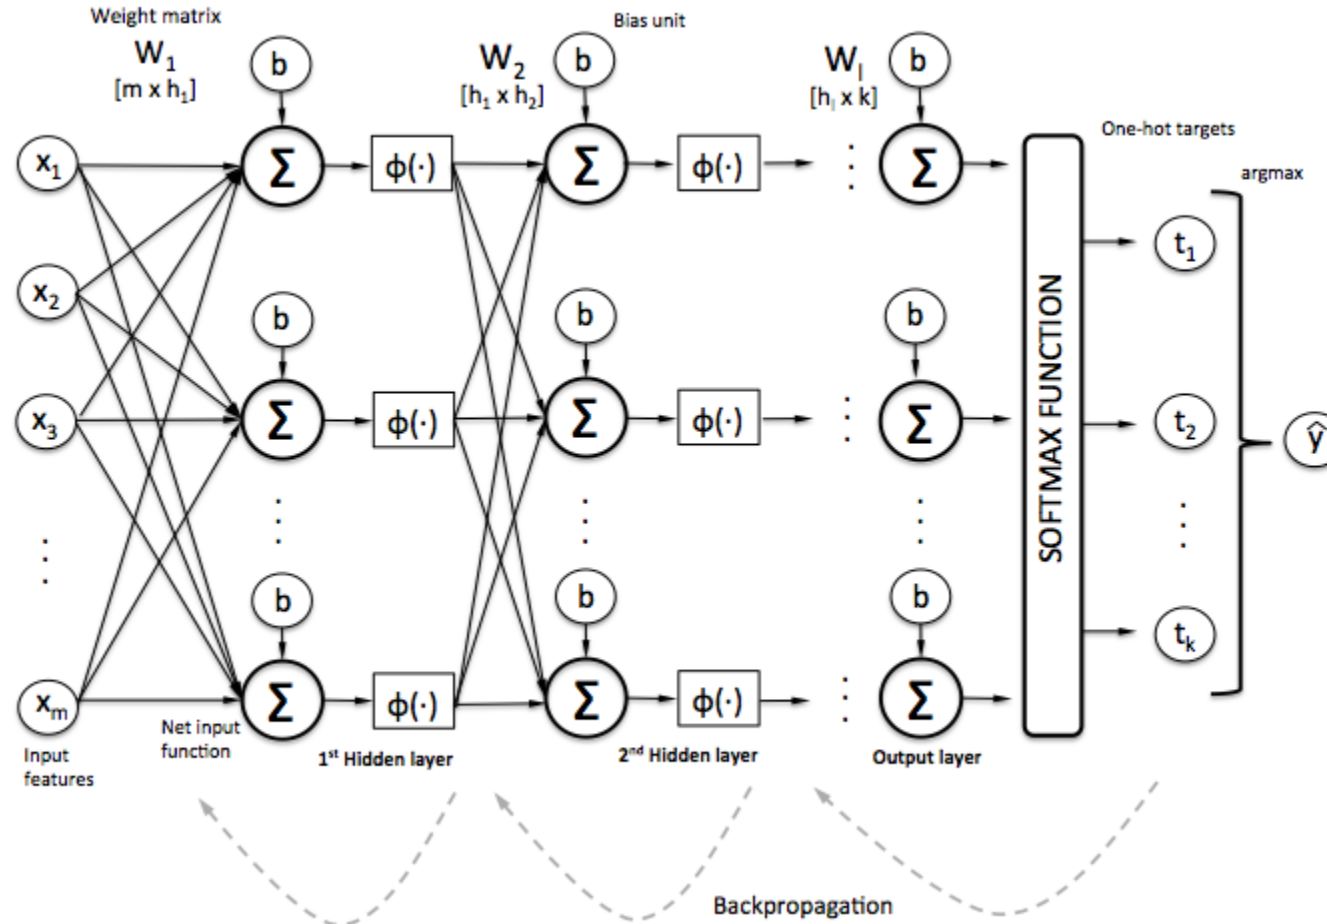
# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Tensors in Flows

- The **tensor** and how does it have 'flow'?
  - **Tensors are everywhere** in AI machine learning mathematic expressions:
    - A vector is a list of values.
    - A matrix is a table (or list of lists).
    - And more:
      - a list of tables (or list of lists of lists)
      - a table of tables (or list of lists of tables...).
      - And so on.
- Let's take a **multi-layer neural network** as an example:
  - Input data features ('x1', 'x2', ...) going through 2 hidden layers:
    - Each with nodes ('neurons')
    - Each with weights ('W') and bias ('b')
  - Output is y.

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Tensors in Flows



Sources: gk\_ at <https://chatbotslife.com/>

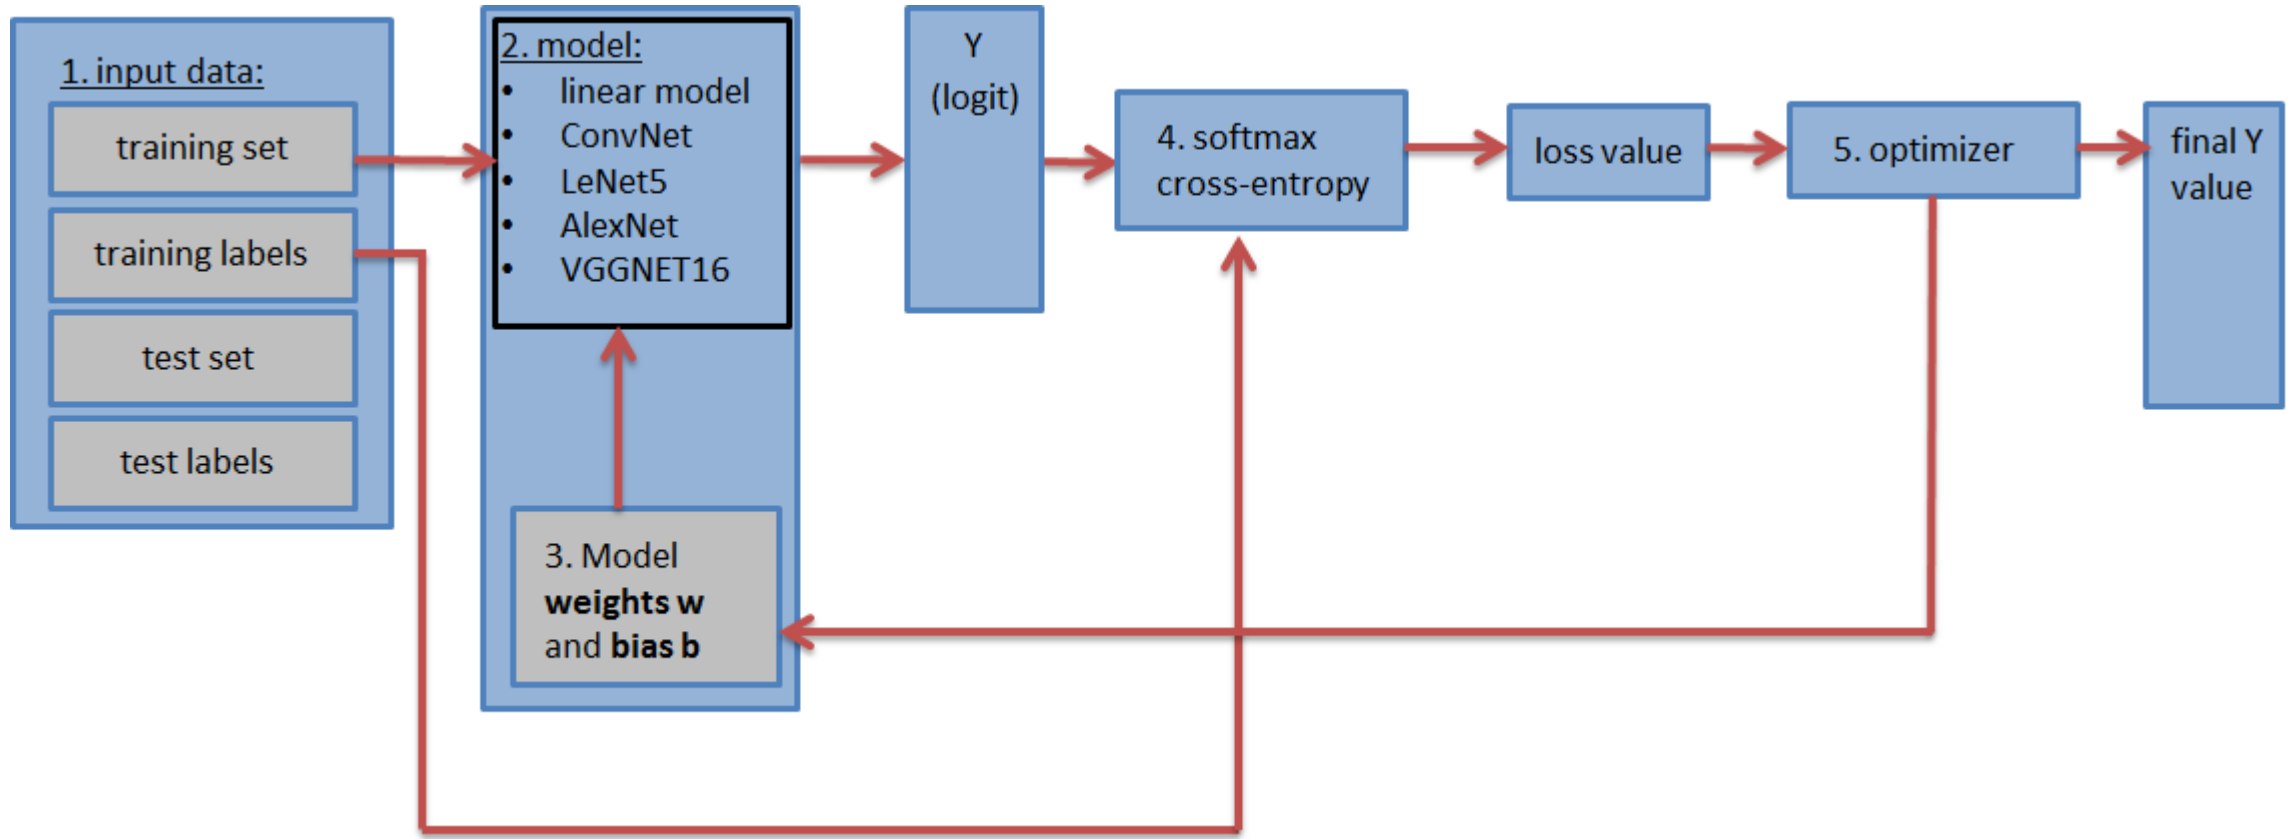
# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Artificial Neural Networks with Tensorflow

- The graph containing the Neural Network should contain the following components
  - The **input datasets**: the training and the test dataset and labels (and the validation dataset and labels).
    - The test and validation datasets can be placed inside a `tf.constant()`.
    - The training dataset is placed in a `tf.placeholder()` so that it can be feeded in batches during the training (stochastic gradient descent).
  - The **artificial neural network model** with all of its layers.
  - The **weight matrices and bias vectors** defined in the proper shape and initialized to their initial values. (One weight matrix and bias vector per layer.)
  - The **softmax cross-entropy and loss value**:
    - The model has as output the logit vector (estimated training labels).
    - By comparing the logit with the actual labels, it is possible to calculate the loss value (with the softmax with cross-entropy function).
    - The loss value is an indication of how close the estimated training labels are to the actual training labels and will be used to update the weight values.
  - An **optimizer**, which will use the calculated loss value to update the weights and biases with backpropagation.

# AI Deep Learning: Introduction to TensorFlow

## TensorFlow: Artificial Neural Networks in Tensorflow



Sources: Ahmet Taspinar